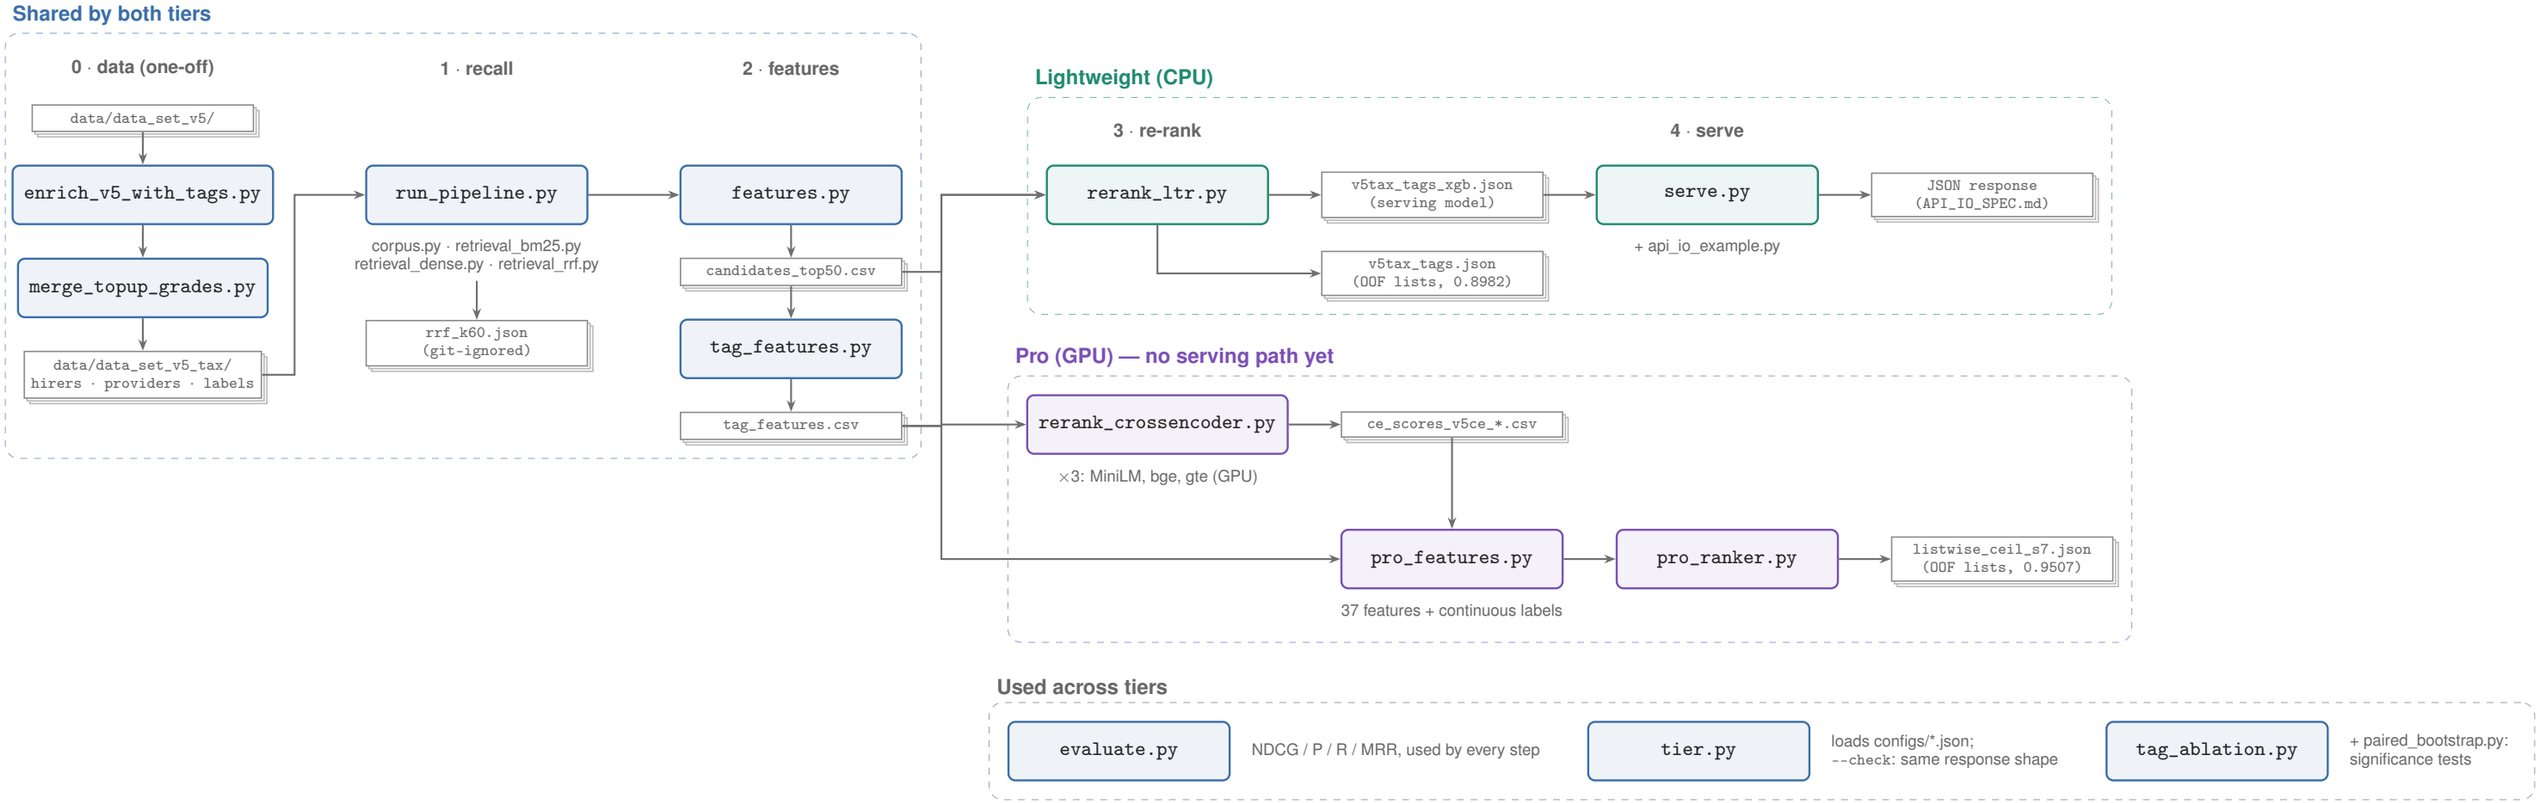

Shared by both tiers
0 · data (one-off)
1 · recall
2 · features
data/data_set_v5/
enrich_v5_with_tags.py
merge_topup_grades.py
data/data_set_v5_tax/
hirers · providers · labels
run_pipeline.py
corpus.py · retrieval_bm25.py
retrieval_dense.py · retrieval_rrf.py
rrf_k60.json
(git-ignored)
features.py
candidates_top50.csv
tag_features.py
tag_features.csv
Lightweight (CPU)
3 · re-rank
4 · serve
rerank_ltr.py
v5tax_tags_xgb.json
(serving model)
v5tax_tags.json
(OOF lists, 0.8982)
serve.py
+ api_io_example.py
JSON response
(API_IO_SPEC.md)
Pro (GPU) — no serving path yet
rerank_crossencoder.py
×3: MiniLM, bge, gte (GPU)
ce_scores_v5ce_*.csv
pro_features.py
37 features + continuous labels
pro_ranker.py
listwise_ceil_s7.json
(OOF lists, 0.9507)
Used across tiers
evaluate.py
NDCG / P / R / MRR, used by every step
tier.py
loads configs/*.json;
--check: same response shape
tag_ablation.py
+ paired_bootstrap.py:
significance tests